



同濟大學
TONGJI UNIVERSITY



Politecnico
di Torino

Tongji University

MAURIZIO VASSALLO

Student ID:

Academic Tutor: a

A Collision Avoidance System Using Reinforcement Learning

Contents

1	Introduction	2
2	Reinforcement Learning	4
2.1	Introduction	4
2.2	Reinforcement Learning compared to other methods	4
2.3	Application of Reinforcement Learning	5
2.4	Formal definition	6
2.5	Q-Learning	8
3	Project Development	10
3.1	Environment	10
3.1.1	Simple Track	10
3.1.2	Complex Track	10
3.2	Agent	10
3.3	State	11
3.4	Policy	11
3.5	Reward Fuction	12
3.6	Actions	12
4	Conclusion	13

1 Introduction

This report is drawn up after the development of the thesis project at Tongji University (同济大学) for the Double Degree project Politong.

The thesis project aims to develop a machine learning algorithm that allows an autonomous vehicle to avoid obstacles.

The research will focus on Reinforcement Learning methods.

Nowadays it is possible to hear more and more about Autonomous Driving Vehicles. The autonomous cars(also known as a self-driving cars or a driverless cars) are a vehicle that are capable of sensing their environment and to navigate without human input. The ability of autonomous vehicles to operate without human intervention depends on their level of technological sophistication, in accordance with the current six-degree autonomy scale proposed by the International Society of Automotive Engineers (**SAE**)[2]. There are 6 levels of driving automation, from level 0 (no automation) to level 5 (full unlimited automation); intermediate levels (1 to 3) are considered semi-autonomous[3, 4]. Currently we are between level 2 and 3, so even if the current technology is behind the famous Level 5 of driving automation, there is lot of work to make it happen.

According to research firm, autonomous vehicles will match or exceed human safety by late 2020s and fulfill all mobility needs in 2040s to 2060s. Optimists predict that by 2030, autonomous vehicles will be sufficiently reliable, affordable and common to displace most human driving, providing huge savings and benefits. However, there are good reasons to be skeptical. There is considerable uncertainty concerning autonomous vehicle development, benefits and costs, travel impacts, and consumer demand. Considerable progress is needed before autonomous vehicles can operate reliably in mixed urban traffic, heavy rain and snow, unpaved and unmapped roads, and where wireless access is unreliable[1].

The idea behind these cars is quite simple: outfit the vehicles with sensors that can track all the objects nearby and make the cars understand the world around them. Autonomous vehicles are driven using technology such as GPS, odometry, radars, laser lights and other devices[4]. These sensors themselves do not make the car ‘smart’, what make it autonomous are the big computers inside and the algorithms they are running. Usually these softwares run neural networks: these take as input the sensor recordings, elaborate them and output some values like steering angle, accelerate, brake or other important values. Even if the idea behind these technological innovative vehicles is simple the implementation is not so straightforward, for some reasons: not enough hardware computation, not enough training data, problems with handle different weather conditions (fog, rain, snow, etc.), the current regulation remains in a nascent stage.

Even if this technology is not diffused yet, there are many potential benefits that autonomous vehicles could introduce in our society:

- **Transportation Safety.** The most notable predicted benefit of autonomous vehicle technology is a substantial reduction in the human and economic toll of traffic accidents. Indeed, impairment, distractions, and fatigue alone account for over 50% of all fatal crashes. The use of autonomous vehicles could significantly reduce the incidence of such crashes, as vehicles with no-human operators are never drunk, distracted, fatigued, or otherwise susceptible to human failings.
- **Access to Transportation.** Another important potential benefit of autonomous vehicle technology is increased mobility for populations currently unable or not permitted to operate traditional vehicles. These populations include older citizens, the disabled, people too young to drive, and others without a driver's license.
- **Traffic Congestion and Land Use.** Autonomous vehicles could reduce congestion and change the way in which cities are planned. Most cars are moving only for 5% of their lives, for the 95% they are parked[5]! For this reason a lot of space is dedicated to parking lots; the same space that could be used for different purposes (green spaces, etc.).
- **Energy and Emissions.** Autonomous vehicle technology has the potential to reduce both energy consumption and pollution thanks to efficiencies gained through smoother acceleration and deceleration and increased roadway capacity.[6]

To do: find a way to link this part with the following one!!

2 Reinforcement Learning

2.1 Introduction

Humans and animals learn through a process of trial and error. This process is based on a reward mechanism that provide a response to our behaviors. The goal of this process is to incentivize the repetitions of actions which trigger positive rewards and disincentivize the repetition of actions which trigger negative ones. Inspired by how animals and humans learn, **Reinforcement Learning** is built around the idea trial and error from an interaction with the enviroment

Reinforcement Learning tries to solve the problem in which a decision-maker, Agent, can perform some action inside a world, called Enviroment. The agent senses the enviroment through the state; for each state the agent has to perform an action. These actions result in an effect: they change the agent's state and give it a feedback, the Reward. The reward is a value which indicated if the action performd is good, then the reward is positive, or it is bad, the reward is negative. Given these rewards the agent understand what action to perform in a given state to get a positive reward.

This cycle of state-action-reward is repeted many times. The goal of the agent is to find a policy such that the actions perform lead to the maximum reward possible. In partical the agent wants to maximize the cumulative reward, the sum of rewards over a period of time. The difference between maximizing the instant reward and cumulative reward is an important distinction because: choosing an action which lead to the maximum next reward could in the next state bring to an end state (game over); instead maximizing the cumulative reward means thinking in a long-term horizon, so that the next states should all have a positive reward (even if the action chosen in one state is not the one that the maximum reward).

2.2 Reinforcement Learning compared to other methods

Machine learning is the science of getting computers to act without being explicitly programmed. Machine Learning methods are appropriate in application settings where people are unable to provide precise specifications for desired program behavior, but where examples of desired behavior are available, or where it is possible to assign a measure of goodness to examples of behavior[7].

There are mainly 3 methods to train Machine Learning algorithms, each with its advantages and disadvantages. The 3 methods are:

- Supervised Learning. In this method the algorithm is trained with labeled data; the training examples are of the form (x_i, y_i) , with x_i a n -dimensional vector and y_i a scalar (it can represent either a class or a floating value).

The learner (either a classifier or a predictor) tries to find a good mapping function that maps an input x_i to its corresponding y_i . During this learning process the error between the y_i predicted and the actual value is calculated and used to make the method learning (usually with Gradient Descent and Back Propagation) and decreasing the error over time. A popular example is to classify an image given the raw pixels.

- **Unsupervised Learning.** The main difference between Unsupervised Learning and Supervised Learning is that in the former method data does not have labels. In this setting, the goal is usually find the relationship in the elements of the dataset. This is done calculating the distance (Euclidian, Hamming, etc.) between the points: if the distance between the points is small, they may share similar characteristics. A popular example is detecting a person purchase preferences analyzing his shopping list with other people shopping's lists.
- **Reinforcement Learning.** Reinforcement Learning comes into play when examples of desired behavior are not available but where it is possible to score examples of behavior according to some performance criterion[7]. In general, in Reinforcement Learning the goal is to maximize an unknown reward function through a trial and error process. A popular example is a car that learns to drive in a given environment with no prior knowledge given.

Supervised and Reinforcement Learning could be considered similar but there are few key differences, in particular the goal in Supervised Learning is to predict the right label while in Reinforcement Learning the goal is to find an action x^* in order to maximize the reward. For this reason instead of minimizing the error between the predicted class and the real class, the goal is to choose an action that maximize the cumulative reward in a given state.

2.3 Application of Reinforcement Learning

Some possible application of RL:

- **Self-driving cars.** Train a car in a real world can be dangerous (people can be hurt) and expensive (car can hit an object and damage the vehicle). They can be trained in a RL setting: the car is the agent, the world around it the environment, it has to take some actions (throttle, steer, etc.), the state are the sensors and the goal is to reach a destination avoiding obstacles. In this setting the agent can get positive or negative rewards depending on the action chosen.
- **Games.** Popular examples are: chess, Go and Dota. In all these games computers were able to beat the champions.

- Finance. Many problems can be formulated as a RL problem, for example the stock trading. Here the action could be: selling, buying, holding and the goal is to maximize the cumulative return over a period of time.
- many other applications[8].

2.4 Formal definition

The Reinforcement Learning can be formulated as a Markov Decision Process (MDP), indeed a MDP express the problem of sequential decision-making, where for each state s the decision maker can choose any action a available in that state s . The process respond moving with some probability to the state s' and giving the decision maker a reward $R_a(s, s')$ (read as: ‘the reward given when in state s and the action a chosen brings to the next state s' ’)

The MDP is defined as a tuple of 4 elements (S, A, P, R), where:

- S is a set of states, called the *state space*.
- A is a set of actions, called the *action space*.
- P is the probability from state s , at time t , of reaching state s' , at time $t + 1$ with action a :

$$P_a(s, s') = Pr(s_{t+1} = s' | s_t = s, a_t = a)$$

- $R_a(s, s')$ is the immediate reward received after transitioning from state s to state s' , due to action a .

The state and action spaces may be finite or infinite.

The MDP is controlled by a sequence of discrete time steps that create a trajectory v :

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$

where the states follow the state transition $P_a(s, s')$. The transition function and the reward function are determined only by the current state, and not from the previous states. This property is called Markov property, which is a characterize the MDP and it means that the process is memory-less and that the future states depends only on the current one and not on its history.

The goal of the MDP is to find a good policy for the decision maker: a function π that specifies the action $\pi(s)$ that will be chosen when in state s . The policy π found will maximize the cumulative reward over a trajectory v :

$$G(v) = \sum_{t=0}^{\infty} R_{a_t}(s_t, s_{t+1})$$

This return value has the problem that all the rewards contribute in the same weight and this can create some problems due to the lack of information. A better return value would be to give some more importance to the short-term memories and giving less importance to the one far in the future. This is solved introducing a **discount factor**, denoted with γ . Then the corrected formula is:

$$G(v) = \sum_{t=0}^{\infty} \gamma^t R_{a_t}(s_t, s_{t+1})$$

with value of γ satisfying $0 \leq \gamma \leq 1$. A lower discount factor motivates the decision maker to take actions that are close in time instead of actions that are far in the future.

<http://incompleteideas.net/book/ebook/node34.html>

<https://cims.nyu.edu/~donev/Teaching/WrittenOral/Projects/XintianHan-WrittenAndOral.pdf>

Another important notion in MDP and Reinforcement Learning is the value function. While the return $G(v)$ gives the reward over a trajectory it does not tell much about how good are the single states. The value function does exactly this, it estimates how good is for the decision maker to be in a given state. The notion of "how good" is defined in terms of future rewards that can the decision maker expect in terms of expected return.

the value function $V_{\pi}(s)$ can be formally defined as:

$$V_{\pi}(s) = \mathbb{E}_{\pi}(R_t | s_t = s) = \mathbb{E}\left(\sum_{t=0}^{\infty} \gamma^t R_{a_t}(s_t, s_{t+1}) | s_t = s\right)$$

The expected return when starting at state s and following policy π .

Similarly, another notion: the action-value function, the expected return from state s with an initial action a :

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}(R_t | s_t = s, a_t = a) = \mathbb{E}\left(\sum_{t=0}^{\infty} \gamma^t R_{a_t}(s_t, s_{t+1}) | s_t = s, a_t = a\right)$$

The value function and the action-value function are called V-function and Q-function. They are related and satisfy a particular relationship, used in many Reinforcement Learning contexts, that for any policy π and state s , the following condition holds:

$$V_{\pi}(s) = \mathbb{E}_{\pi}(Q_{\pi}(s, a))$$

The V-function can be decomposed in 2 terms:

$$V_{\pi}(s) = \mathbb{E}_{\pi}(R_t | s_t = s) = \underbrace{\mathbb{E}_{\pi}(R_t | s_t = s)}_{\text{immediate reward}} + \underbrace{\mathbb{E}_{\pi}(\gamma V_{\pi}(s_{t+1}) | s_t = s)}_{\text{discounted value of next state}} \quad (1)$$

This is the Bellman Equation (1) that defines the value function recursively, enabling the estimations of the next states. Similarly it is possible to write Bellman equation for the Q-function :

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}(R_t | s_t = s, a_t = a) = \mathbb{E}_{\pi}[R_t + \gamma Q_{\pi}(s_{t+1}, a_{t+1}) | s_t = s, a_t = a]$$

In this way the V-function and the Q-function are updated with the values of the successive states without the need to know the trajectory till the end.

2.5 Q-Learning

Q-learning is an off-policy reinforcement learning algorithm that tries to find the best action to take given the current state. It's considered off-policy because the q-learning function learns from actions that are outside the current policy, like taking random actions. Off-policy means that the update can be made independently from the policy that has gathered the experience. This means that off-policy algorithms can use experiences collected in the past to improve the policy.

Q-Learning is also a TD algorithm (Temporal Difference algorithm[9]) and it inherits from the TD learning the characteristics of one-step learning.

These two elements (TD and off-policy property) are important; indeed the goal of Q-learning is to approximate the Q-function using the action that maximizes the Q-value of the next state. Formally:

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha[r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)] \quad (2)$$

where α is the learning rate, that indicates how much the Q-value will be updated, and γ the discount factor, that gives more importance to the short-term actions and gives less importance to the long-term ones.

This modified version of the Bellman equation updates the value of the current Q-value using only the next step (using the TD one-step learning property) and it is an off-policy algorithm since it uses the \max_a to choose a_{t+1} .

Simple Q-learning problems can be solved using a Q-table, where for each state there are all the Q-values for all possible actions:

States	Actions			
	a_1	a_2	\dots	a_m
s_1	$Q(s_1, a_1)$	$Q(s_1, a_2)$	\dots	$Q(s_1, a_m)$
s_2	$Q(s_2, a_1)$	$Q(s_2, a_2)$	\dots	$Q(s_2, a_m)$
\dots	\dots	\dots	\dots	\dots
s_n	$Q(s_n, a_1)$	$Q(s_n, a_2)$	\dots	$Q(s_n, a_m)$

Table 1: state-action Q-values

with s_i a t -dimensional array (imagine a possible car's state which at each instant record t values from different sensors). During training the Q-values are updated using equation the Bellman equation for Q-learning (2).

This implementation has 2 main problems:

- It is possible to see that the size of the table increases exponentially with the increasing number of states (n) and the number of actions (m), increasing the memory required to store the table.
- The amount of time required to explore each state to create the required Q-table would be increase too.

One way to solve this problem is to use a neural network. This new architecture assumes the name of **Deep Q-learning**. Deep Q-learning replaces the Q-table with a neural network that approximate the Q-value fuction; moreover rather than mapping a state-action pair to a Q-value, a neural netowk maps an input state to an $(action, Q-value)$.

This solution overcomes the 2 problems mentioned above.

In this contest, the objective is to change the neural network's weights θ , so that Q_θ is an optimal Q-value fuction. In order to change the weights, as common with neural networks, the Gradient Descend is employed. In particular the weights are changed with the following formula:

$$\theta = \theta - \alpha \nabla_{\theta} \mathbb{E}_{(s,a,r,s')} [(Q_{\theta}(s_i, a_i) - \text{target}_i)^2]$$

where:

$$Q_{\theta}(s_i, a_i)$$

is the value the neural network output;

$$\text{target}_i = r_i + \gamma \max_{a'_i} Q_{\theta}(s'_i, a'_i)$$

is the value the neural network should output;

$$\mathbb{E}_{(s,a,r,s')} [(Q_{\theta}(s_i, a_i) - \text{target}_i)^2]$$

is the Mean Square Error over a batch iteration;

$$\nabla_{\theta}$$

is the gradient of the loss w.r.t the neural network's weights.

3 Project Development

The goal of the project is to use some Reinforcement Learning algorithms to teach an actor to avoid obstacles. The **Agent** is put in an **Environment** it does not know and it will learn a good **Policy** in order to avoid the walls using a trial and error approach. The agent will be given with positive or negative **Rewards** depending on the **Actions** it performs.

3.1 Environment

In this project a deterministic environment will be used since the next state of the environment can always be determined based on the current state and the agent's action. In particular there are 2 main environments: A simple track, and a complex one used for testing purposes. Both tracks were modelled in **Blender** version 2.93[10].

These models are used inside of Unreal Engine 4 [11], a powerful game and physics engine.

3.1.1 Simple Track

This simple track is used for training and validation. Since this is used for training it does not have long straightaways or sharp turns. During training the track is run clockwise and counterclockwise so that the agent learns to turn right and left.

3.1.2 Complex Track

This complex track is used for testing. This will have more long straightaways or sharp turns.

3.2 Agent

The agent is a car available on the Vehicle Variety pack [12] asset of Unreal Engine 4. From this packet only the 'SportCar' content is used and some settings are modified (drag coefficient and other physical values).




Figure 1: The used car as Agent with its own measurements

3.3 State

The agent senses the environment around it using some LIDAR sensors (light detection and ranging [13]).

The number of Lidar to use is an important hyper-parameter: in this case good results were found using 32 Lidar sensors. This 32 sensors are equally spaces in an angle range of 170° , this is another important hyper-parameter. Multiple tests were made to set these hypre-parameters, in particular:

- Too many Lidar sensors requires more training time, they increase the performance but sometimes the model did not converge.
- Too few Lidar sensors the algorithm did not have enough information to learn properly.
- Too wide angle decresed the sensor density and decresed the focus on the most important part: the front and the sides of the car.
- Too small angle increased the sensor desity and increased the focus only on some specific region: the front.

So a trade-off between training time and performances were reached with 32 sensors and a 170° angle

3.4 Policy

The policy, as mentioned early, is a mapping from state to action, how the agent chooses an action. The goal is to improve the policy over time in order to

increase the cumulative reward.

As common in Reinforcement Learning, the agent needs to explore the environment to get to know the world around it and after some time then apply this knowledge. This is known as the **exploration-exploitation trade-off**: at the beginning the agent chooses random actions (exploration), after it has learnt it uses what it knows to choose the correct action (exploitation).

This exploration-exploitation trade-off is implemented with a value ϵ commonly initialized to 1: at each step a random value c (with $c \in [0, 1]$) is drawn, if $c \leq \epsilon$ then the agent chooses a random action and ϵ is decreased otherwise it selects the action a with the maximum Q-value given the current state s ($\max_a Q_\theta(s, a)$). The value of ϵ is decreased with a value $\epsilon_{decay} = 7.5 \cdot 10^{-5}$; once ϵ keeps decreasing the probability that $c \leq \epsilon$ will be lower so preferring exploitation over exploration.

The ϵ_{decay} is a hyper-parameter: the lower ϵ_{decay} is the more the agent will explore the environment increasing the training time.

The performance of the policy depends on the neural network structure. In this project 3 different architectures are used. The idea behind them is similar what changes is how the next expected Q-value; the Bellman equation changes.

3.4.1 Deep Q-Network

3.4.2 Double Deep Q-Network

3.4.3 Dueling Double Deep Q-Network

3.5 Reward Function

3.6 Actions

4 Conclusion

This internship was a formative experience that allowed me to discover many new topics but also to understand how a company works. Moreover, it allowed me to understand that there is a substantial difference between the university environment and the corporate environment: university, as mentioned by some professors during the courses, has the task of giving a general knowledge of the topics and the real task is to give a way of thinking and analysing a problem to find a solution; indeed, during this period in the company I realized that computer knowledge was not always essential but it was more important to understand what was needed to be done and how to do it.

I believe that this internship was very useful because it allowed me to deepen and learn about new topics and gave me an idea to understand if the tasks assigned could be addressed in a future university career or in a possible job. This is also an experience that allows to grow both personally and professionally, not to mention that it is a way to enrich the Curriculum Vitae.

I am fully satisfied with the internship as AI researcher and developer and I am satisfied with the atmosphere of serenity and professionalism that I lived in the company thanks to the kind and extremely helpful colleagues.

References

- [1] Todd Litman. 2021. *Autonomous Vehicle Implementation Predictions*.
<https://www.vtapi.org/avip.pdf>
- [2] SAE. 2016. *Taxonomy and Definitions for Terms Related to Driving Automation Systems for On-Road Motor Vehicles*.
https://www.sae.org/standards/content/j3016_201806/
- [3] Monika Stoma, Agnieszka Dudziak, Jacek Caban and Paweł Drozdzie. 2021. *The Future of Autonomous Vehicles in the Opinion of Automotive Market Users*.
<https://www.mdpi.com/1996-1073/14/16/4777/pdf>
- [4] Piotr CZECH, Katarzyna TUROŃ, Jacek BARCIK. 2018. *AUTONOMOUS VEHICLES: BASIC ISSUES*.
<https://doi.org/10.20858/sjsutst.2018.100.2>
- [5] Paul Barter. 2013. *"Cars are parked 95% of the time". Let's check!*.
<https://www.reinventingparking.org/2013/02/cars-are-parked-95-of-time-lets-check.html>
- [6] Jeremy A. Carp. 2018. *AUTONOMOUS VEHICLES: PROBLEMS AND PRINCIPLES FOR FUTURE REGULATION*.
<https://scholarship.law.upenn.edu/cgi/viewcontent.cgi?article=1048&context=jlp>
- [7] ANDREW G. BARTO and THOMAS G. DIETTERICH. *Reinforcement Learning and its Relationship to Supervised Learning*.
<http://web.engr.orst.edu/~tgdp/publications/Barto-Dietterich-03.pdf>
- [8] Yuxi Li. 2019. *REINFORCEMENT LEARNING APPLICATIONS*.
<https://arxiv.org/pdf/1908.06973.pdf>
- [9] Scholarpedia.org *Temporal difference learning*.
http://www.scholarpedia.org/article/Temporal_difference_learning
- [10] *Blender*.
<https://www.blender.org/>
- [11] *Unreal Engine 4*.
<https://www.unrealengine.com/>

- [12] *Vehicle Variety Pack*.
<https://www.unrealengine.com/marketplace/en-US/product/bbcb90a03f844edbb20c8b89ee16ea32>
- [13] *LIDAR*.
<https://en.wikipedia.org/wiki/Lidar>